

Jeet Thakore

jeetpt2@illinois.edu | [linkedin/jeetthakore](https://www.linkedin.com/in/jeetthakore) | jeetthakore.in | [github/JeetThakore](https://github.com/JeetThakore) | +1 (447) 9026-701 | Authorized to work in the U.S

EDUCATION

University of Illinois Urbana-Champaign	August 2024-May 2026
Master of Science in Information Management (STEM) (GPA: 3.81/4.00)	
Institute of Infrastructure Technology Research and Management	August 2018-May 2022
Bachelor of Technology in Electrical Engineering (GPA 3.84/4.00)	

WORK EXPERIENCE

Data Engineer	June 2025-August 2025
Astellas Pharma Inc. (U.S.A)	
- Engineered 6 Azure Data Factory pipelines landing clinical trial data from SFTP and REST APIs into ADLS Gen2, then converted 40 GB of raw healthcare data to partitioned Parquet with PySpark on Azure Databricks, cutting Spark SQL scan time from 11 minutes to under 2 minutes - Wrote SQL reconciliation checks comparing source control totals against curated clinical data, surfacing a truncation bug where the legacy loader silently cut fields over 255 characters, and exposed 9 curated tables via serverless SQL pools in Azure Synapse Analytics - Extended Terraform IaC modules to provision storage containers and Databricks cluster policies scoped to HIPAA compliance controls on patient health information (PHI), replacing a manual runbook, and wired it into Azure DevOps (YAML CI/CD) with Git branch validation	
Data Engineer: Assistant Systems Engineer	July 2022-May 2024
Tata Consultancy Services Limited	
- Implemented SCD Type 2 on counterparty, instrument, and book dimensions using Delta Lake MERGE with hash based change detection in PySpark, replacing full daily reloads and reducing dimension refresh time from 50 minutes to 6 minutes for month end financial reporting - Rebuilt 60 ad hoc Spark SQL views into a dimensional model of 8 conformed dimensions and 4 fact tables on Databricks, eliminating duplicate transformation logic across revenue cycle management and treasury reporting - Resolved an Apache Spark data skew issue where 3 of 200 partitions contained 58% of records using salted join keys and Adaptive Query execution, restoring nightly financial loads to a 4 hour SLA while reducing Databricks compute spend by 34% through cluster optimization - Orchestrated 30+ Azure Data Factory pipelines using metadata driven control tables and implemented Unity Catalog governance, column-level masking on 11 restricted fields, and Azure DevOps CI/CD promotion gates with Git pull request validation	
Data Engineering Co-Op	January 2022-April 2022
Tata Consultancy Services Limited	
- Created Apache Airflow DAGs to orchestrate ingestion of 10M+ EHR records and ICD-10 diagnosis data from MongoDB into Amazon Redshift, automating daily workflows supporting Clinical Data Warehousing and healthcare analytics - Implemented data validation checks, workflow monitoring and retry policies using Airflow and Bash Operators while operating within established Role-Based Access Control (RBAC) and Data Governance controls for clinical datasets	

LEADERSHIP EXPERIENCE

Growth Cohort Lead	January 2026-May 2026
Gies College of Business - University of Illinois Urbana Champaign	
Led performance visibility and accountability for a 180-person portfolio, driving structured progress tracking and evaluation across cohorts - Coached 20 individuals across two cohorts through weekly accountability sessions focused on measurable progress improvement - Owned program operations including onboarding, engagement, and evaluation, ensuring reliable execution across cohorts at scale	

PROJECTS

Containerized Data Pipeline: Real-Time API Streaming with Kafka, Spark, Airflow & PostgreSQL
- Architected a containerized ETL pipeline streaming REST API data through Kafka, transforming with Spark, and loading into PostgreSQL - Configured Airflow DAGs for task scheduling, decoupling ingestion from processing with fault-tolerant Kafka queuing
Real-Time Heart Rate ELT Pipeline: Kafka Streaming to Snowflake with Airbyte, dbt, Dagster & Tableau
- Streamed simulated IoT heart rate data through Confluent Kafka into S3, ingesting incrementally into Snowflake via Airbyte with CDC - Modeled a star schema with SCD Type 2 dimensions in dbt, scheduled daily through Dagster Cloud with auto-materialization policies
UIUC Student Assist: AI Chatbot for Campus Information
- Built scalable data ingestion pipeline processing 1,500+ documents into embeddings, automating chunking and FAISS vector indexing - Integrated LangChain RAG workflow with fine-tuned FLAN-T5, delivering real-time Q&A via Streamlit and reducing search time by 70%

SKILLS

Languages & Big Data: Python, SQL (PostgreSQL, MS-SQL, T-SQL), NoSQL (MongoDB), Java, Apache (PySpark, Spark SQL, Airflow, Kafka)
Cloud Platforms: AWS (Glue, EMR, Redshift), Azure (Data Factory, Synapse, ADLS), BigQuery, Snowflake, Kubernetes, Terraform
Databricks Ecosystem: Medallion Architecture, dbt, Delta Lake, Delta Live Tables, Unity Catalog, Auto Loader, Lakeflow, MLflow
Data Engineering: Data Warehousing, Dimensional Modeling, Data Governance, Data Lineage, Data Validation, CDC, HIPAA, GDPR
AI & Analytics: LangChain, RAG, LLM Fine-Tuning, NLP, Scikit-Learn, Pandas, NumPy, Power BI, Git, CI/CD, Linux

MICROSOFT CERTIFICATIONS

- Microsoft Certified: Azure Data Scientist Associate (DP-100)
- Microsoft Certified: Power BI Data Analyst Associate (PL-300)